

System Design :

Applying CS concepts (networks, distributed systems, parallel computing) to build large scale systems.

- 1) High Availability requests should not fail
- 2) Scalability handle increasing workload
- 3) Consistency same information across all servers
- 4) Fault Tolerance function despite component failures
- 5) Load Distribution evenly spread traffic / workload
- 6) Extensibility easy to add / modify components
- 7) Performance speed and efficiency

Primers —

- Vertical Scaling
- CRON jobs
- Resilience and Backups (Master-Slave Architecture)
- Horizontal Scaling
- Microservices
- Distributed Systems
- Load Balancing
- Decoupling
- Logging and Metrics (Analytics, Auditing, Reporting, ML)
- Extensibility
- High-level Design
- Low-level Design

Code on computer = takes input → output

To share: expose code via API

Request → input sent by user

Response → output returned by system

Problems with self hosting

- Requires database setup and endpoint configuration
- Risk of downtime

Solutions is Cloud

- Cloud = provider's computers rented (e.g. AWS)
- Offers computing power, reliability, configuration management
- Remote login allows hosting services on cloud

Scaling for growth

1) Vertical Scaling → bigger machines (more CPU, RAM)

2) Horizontal Scaling → more machines, load distributed

.. Vertical Scaling

- Fast (interprocess communication)
- Consistent data (single source of truth)
- Single point of failure
- Hardware limit
- Load balancing not required

.. Horizontal Scaling

- Slow (network calls, remote procedure calls)
- Data consistency issues (need loose guarantees)
- Resilient (failure handled by other machines)
- Scales well (linear with more servers)
- Load balancing required

Real World Solution (Scalability) - Hybrid Approach

- Start with vertical scaling
- Move to horizontal scaling as users grow
- Vertical → speed + consistency
Horizontal → scalability + resilience

Server - Client :

- Server = processes requests
- Client = sends requests
- Response = processed result

Scaling Problem :

- 1) One server → fine for small load
- 2) Many clients → multiple servers needed

Load Balancing

- Evenly distribute requests across N servers.
- Simple method : hash (request ID) % N
- X requests → X/N load → Load factor : 1/N
- Works well initially.

Problem with Naive Hashing

- Adding / removing a server changes N.
- Causes remapping of almost all requests → cache invalidation & instability
- Ex : Going from 4 servers to 5 → nearly 100% requests remapped
- Request stickiness : same user ID should map to same server for caching
- Naive re-hashing → breaks cache locality → wasted computations and poor performance.

Consistent Hashing

It ensures that when servers are added / removed, only a small fraction of requests are remapped, preserving cache locality and balancing load efficiently.

Instead of remapping all requests:

- 1) only a small portion of requests get redistributed when servers are added / removed.
- 2) Each server owns a segment of a hash ring ($0 \dots M-1$)
- 3) Adding a new server = steals small sliver from all existing servers (not full redistribution).

How it works?

- Hash Requests \rightarrow map request IDs to positions on ring
- Hash Servers \rightarrow map server IDs (via hash $0 \dots M$) to ring positions
- Request Mapping Rule \rightarrow each request is served by the nearest server clockwise on the ring

Practical Problem

- With few servers \rightarrow skewed distribution possible
ex: one server may take $\sim 50\%$ load if unlucky placement
- Theoretical avg = $1/M$, but variance is high.

Solution \rightarrow Virtual Servers

- Assign multiple points per server on the ring.
- Implementation \rightarrow use K different hash functions per server.
- Each server now appears multiple times \rightarrow load spreads out.
- If $K \approx \log M$, load distribution \approx uniform
- Adding / removing a server \rightarrow only fraction of requests move
- No single server ends up overloaded.

Consistent Hashing + Virtual Servers = Scalable, fault tolerant, and evenly distributed request mapping with minimal remapping.

Messaging Queues / Task Queues

- Producers (clients) send tasks / requests.
- Consumers (servers / workers) process them.
- Requests are placed asynchronously → producer gets an acknowledgement, not the result immediately.
- Tasks go into a queue and are processed in order or based on priority.

Benefits of Asynchronous Processing

- 1} Clients don't wait → can do other tasks
- 2} Server can prioritize tasks
- 3} Better resource utilization + smoother client experience.

Scaling

- Multiple servers can pull from queue
- If one server fails:
 - Non persistent tasks are lost
 - Persistent tasks can be redistributed to healthy servers.

Fault Tolerance

- Simple in-memory queues fail if a server crashes → persistence is required
- Queue systems store metadata in Database such as:
 - Task ID
 - Payload
 - Status (pending, in-progress, completed)

- Mechanism : checks if server alive (e.g. ping every 15s)
 - If dead → redistribute unfinished orders to other servers

Problem → Duplication

- Risk of multiple servers picking same task.
- Solution → Load Balancing (e.g. Consistent Hashing) ensures:
 - ▷ Balanced load
 - ▷ No duplicate requests processed.

Message Queue = Encapsulation of Complexity

Features combined into one system:

- 1) Task persistence (Database)
- 2) Assignment to correct servers
- 3) Heartbeat monitoring (mechanism done by notifier)
- 4) Load balancing + retry if server fails
- 5) Duplicate avoidance

Message Queue is a fault tolerant, asynchronous, persistent, load balanced task handler that encapsulates all the complexity of distributed request processing.

Ex:-

- RabbitMQ
- ZeroMQ
- JMS (Java Messaging Service)
- AWS SQS / SNS

Monolith vs Microservices Misconceptions

- Monolith ≠ one big machine → can scale horizontally across multiple servers + databases
- Microservice ≠ tiny code → it's a business unit (data + functions) bundled together
- Microservices may be few (e.g. 3-5), not "hundreds of tiny ones".
- Clients usually talk via an API Gateway, not directly to services.

Monolith :

Advantages -

- 1) small / cohesive Teams → simpler to manage
- 2) less moving parts → easier deployments
- 3) shared setup code (connections, configs, tests) → no duplication
- 4) Performance → in-process calls (no network latency)

Disadvantages -

- 1) steep learning curve for new devs (need full context)
- 2) coupled deployments → any change = redeploy entire system
- 3) complex testing → everything touches everything
- 4) single point of failure → crash = whole system down

MicroServices :

Advantages -

1. Scalability → scale services independently
2. Easier onboarding → devs only need service context
3. Parallel development → teams work independently
4. Targeted scaling → allocate resources to hot services only.

Disadvantages -

1. complex design → needs strong architecture
2. Over-splitting risk → if two services only talk to each other, they should probably be one
3. RPC overhead → network calls slower vs local calls.

Monolith = simpler, faster, good for small teams (eg StackOverflow)

Microservices = scalable, modular, fault-tolerant, but need smart design (eg. Google FB)

Cache

Caching = storing data temporarily in fast-access memory to avoid repeat computations / queries

Goal → Reduce latency ($\frac{\text{fast}}{\text{response time}}$)

Solve repeated computation or DB lookups (load on DB)

Improve performance, user experience and scalability

Feed Fetching Flow

- Without cache : User → Server → DB query → Server → User
- With cache : similar queries (cohorts of users) can be served from cache instead of querying DB each time
- Result : Latency drops from 220 ms to 202 ms ~10% saving

Why Cache ?

- Reduce latency : Faster retrieval compared to recomputing or fetching from DB
- Reduce load on databases / servers : serve repeat requests from memory
- Improve user experience : Faster responses
- Enable scalability : same system can handle more requests

Cache Placement

- 1) Client-side cache (e.g. app/browser stores data locally)
 - ~~Save~~ network calls
- 2) Server in-memory cache (within the application)
 - Fastest, lives alongside app (e.g. HashMap in RAM)
 - But limited to a single server
- 3) Database cache (built into DB)
 - Common queries auto-cached internally
 - Usually opaque to developers
- 4) Distributed / Global cache (dedicated caching system)
 - Independent service (e.g. Redis, Memcached)
 - Shared across multiple services
 - Can scale independently; update without redeploying main app

Large-scale systems usually use all of them together.

Cache Policies

1. Write Policies (when to update cache vs DB)
 - Write-through : Update DB + cache simultaneously
 - Write-back/behind : Update cache immediately, write to DB later
 - Write-around : Write only to DB, load into cache only on reads
2. Eviction Policies (what to remove when cache is full)
 - LRU (Least Recently Used) - remove oldest accessed item
 - LFU (Least Frequently Used) - remove least accessed item
 - FIFO (First In First Out) - remove oldest inserted item
 - Random / ML based - advanced adaptive methods

Limitations & Trade offs

1) Size Limits

- cannot fit entire DB in memory
- Must choose hot / frequently accessed data

2) Cache Miss & Poor Hit Rate

- If queries don't match cached data → fallback to DB
- leads to wasted lookup + DB call → overhead

3) Threshing

- Frequent evictions and reloads
- causes higher latency + wasted memory usage

4) Stale Data / Eventual Consistency

- Cache is copy of DB → may not always be updated in real time
→ depends on cache update strategy
- Acceptable for views / like, risky / dangerous for critical data

Single Point of Failure :

SPOF = A single component whose failure crashes the entire system

Ex: One database server → if it fails, whole system fails.

Tests resiliency of architecture

Non resilient system = multiple services relying on one fragile point

Mitigation Strategies :

1) Add Redundancy / More Nodes

- Duplicate servers (profile servers, APIs)
- Not useful as "cold backups" unless data is replicated

2) Database Replication (Master - Slave)

- Writes → Master
- Reads → Slaves
- failover → promote slave if master fails
- Failure probability reduces from $P \rightarrow P^2$

3) Load Balancers

- Distribute requests
- But load balancer itself = SPOF
- Solution → Multiple load balancers behind SPOF DNS

4) DNS with Multiple IPs

- Clients resolve same host (e.g. facebook.com) to multiple IPs (different load balancers)

5) Multi Region Deployment

- Avoid full system failure from regional disasters
- Active - active / Active - passive regions

6) Coordinator / Leader Election

- Avoid SPOF in coordinators using similar replication / redundancy strategies.

Ex: Netflix chaos monkey → randomly kills nodes in production to test resiliency
Build pipeline of distributed, fault tolerant system

Content Delivery Network %

CDN = Geo-distributed cache servers that store static content (HTML, videos, images, files) close to users

Purpose → Faster, cheaper, and regulation-compliant content delivery

Why CDN is needed → Central Cache Limitations

- High latency for users far from server (e.g. India vs Japan vs US)
- No single server location suits all → always slow for some region
- Local data regulations (e.g. movies allowed only in India)

Working

- Content is cached in local edge servers around globe
- Users connect to the nearest CDN node instead of central server
- CDN servers = real servers with file system, managed by origin server.

Benefits

1. Low latency → faster webpage loads → higher trust & professionalism - (Amazon/ Google)
2. Compliance → Server region - specific content per regulations
3. Relevance → local caches can store popular / regional data only

CDNs Providers

- Large cloud companies (hard to build your own)
- Ex : AWS CloudFront
 - cheap, reliable, integrated with S3
 - auto-triggered updates when new files added
- Ex : CCP, Azure CDNs

Event-Driven Services

Problem with Request-Response in Microservices

- In microservices, one service often needs to trigger actions in multiple others.
- 1) Services wait for each other → increased latency
 - 2) If one service fails → cascading timeouts
 - 3) Can lead to stale or duplicate data updates
 - 4) Hard to scale when new services need the same events.

Publisher - Subscriber Model

- Services publish events to a Message Broker (Kafka, RabbitMQ).
- Other services subscribe to the events they care about.
- Broker delivers to subscribers asynchronously.
- Fire-and-forget : Publisher doesn't care who consumes, Broker guarantees delivery.

Flow :

Client → Service → Broker → Subscribers

Advantages

1. Decoupling → Services don't know subscribers.
2. Reliability → Broker persists and replays if subscriber down.
3. Faier Dev → One generic interface for publishers.
4. Transaction guarantee → at least once delivery.
5. Scalability → Easy to add new consumers.
6. Simplified architecture → Fewer direct dependencies.

Drawbacks

- ▷ No Strong Consistency / Atomicity
 - Problematic for financial systems as they require atomicity.
 - Cross service transactions can leave inconsistent balances.
- ▷ Idempotency Issues
 - Same message replayed → duplicate operations.
 - Needs Request ID + app logic to enforce idempotency.
- ▷ Not fit for mission critical systems needing strict success/failure guarantees.

Use cases

- Good : Analytics, gaming events, social feeds (Twitter)
- Not good for : Banking, fund transfer, payments

Request Response → pull based

- A service asks another for something, waits for the response
- Tight coupling, dependencies, timeout issues, stale data risk

Event Driven Architecture → push based

- A service publishes an event : "something happened"
- Subscribers decide if it matters to them
- Loosely coupled, scalable, more resilient
- Flow: Producer → Event Bus → Consumers

Ex:

Cli → commits as events

React / Node.js → UI & server-side event driven framework

Hammer → headshots validated by replaying events with timestamps

Advantages

1. Decoupling - services don't depend directly on each other
2. Availability - services persist relevant events locally
system still works even if others are down
3. Event Log (History) - allows debugging, replay, and state restoration
4. Smooth Replacements - new service can replay past events
to sync with old state
5. Transactional Guarantees -
 - At most once (optional, e.g. welcome email)
 - At least once (critical, e.g. invoice email)
6. Scalable - easy to add new consumers without changing producers
7. Storing Intent - keep why a change happened, useful
for future services.

Drawbacks

1. Consistency Issues - replicated local data may get outdated
2. External calls problem - replaying past events may yield different results
3. Lack of fine control - can't easily enforce timing, order, or exact consumers
4. Security concerns - hard to limit who consumes events
5. Event Replay Complexity:
 - Replay from start (impractical at scale)
 - Diff-based (grose changes only)
 - Undo logs (limited, not always possible)
 - Compaction (squash old events into snapshots)
6. Hard to Reason About Flow - tracing event propagation is difficult for developers
7. Vendor Lock-In - moving back to request-response isn't smooth

.. SQL (Relational DBs)

- Data Model → Tables, rows, columns
- Relationships → foreign keys, joins
- Schema → Fixed (rigid, hard to change)
- consistency → strong → ACID (Atomicity, Consistency, Isolation, Durability)
- Best For → financial systems, transactions, structured data, strong consistency
- Ex: MySQL, PostgreSQL, Oracle

NoSQL (Non-Relational DBs)

- Data Model → Document (JSON), Key-Value, Wid-Column, Graph
- Relationships → Not implicit - joins hard (expensive)
- Schema → Flexible (can add/remove attributes anytime)
- Consistency → Eventual (sacrifices strict ACID)
- Best for → Scalability, flexible schema, high write throughput, analytics, caching
- Ex: Cassandra, MongoDB, DynamoDB, Redis

Advantages

- 1) Single Blob Access → Insert/retrieve whole object at once
→ Great for INSERT + SELECT type use cases
- 2) Flexible Schema → Easy to add/remove fields (no costly table ALTER).
- 3) Horizontal Scalability → Inbuilt sharding/partitioning
→ Handles huge scale & availability
- 4) Aggregation Friendly → Optimized for metrics and analytics

Disadvantages

1. Weak Consistency → No guaranteed ACID → bad for transactions ($UPDATE = DELETE + INSERT$)
2. Slow for reads → Reads entire blob even if only 1 column needed
3. No Implicit Relations → Hard to enforce foreign keys
4. Joins are Manual → Expensive and error prone

Use :

- Large scale apps needing availability and writes-heavy workloads
- Flexible schemas, fast evolving products
- Analytics, aggregations, redundancy

Cassandra Architecture

1. Cluster of Nodes → Data spread across multiple machines
2. Hashing (consistent Hashing) →
Keys → Hashed → Assigned to nodes
 - Good hash → equal load distribution
 - Bad hash → imbalance → solved via multi level sharding
3. Replication → Each piece of data stored in multiple nodes
 - Improves availability + fault tolerance
4. Load Balancing → queries routed evenly across nodes
5. High Availability > Strong consistency

Problem : Replicas may not be in sync

- Replication Factor (RF) → no. of copies of data stored
- Quorum → majority must agree on a value (latest timestamp wins)
 - (e.g. RF=3, Quorum = 2 → 2/3 nodes must agree on value)
 - Ensures eventual consistency
 - Accepts small risks of stale reads for high availability
- Distributed consensus → Nodes agree on latest value

Methods : Quorum, Gossip protocol, Paxos, etc.

Data storage

- Writes go to in-memory structure (MemTable)
- Writes stored in log-structured merge (LSM) style
- First in memory (memtable) → flushed to disk as SSTables (Sorted Sorted String Tables)
- Periodic Compaction → Merges SSTables → improves read efficiency
- Benefits
 - Writes are fast (sequential log append)
 - Reads slightly slower (need to check multiple SSTables)

Application Programmable Interface (API)

- A contract that defines how to interact with code, not how it works
- Defines:
 - Function name → what it does
 - Parameters → what it needs
 - Response object → what it returns
 - Possible errors
- Ex: getAdmins (groupId) → returns list of admins (JSON object)

Best Practices

1. Placement → belongs in right service (e.g. group service)
2. Naming → must match the action (getAdmins should return only admins)
3. Parameters → keep minimal, only what's necessary.
extra params only if required for optimization

- 4. Responses → Return only what's needed
· don't overload with unnecessary info
- 5. Errors → Define expected errors (e.g. grp not found)
· avoid overly generic or overly detailed error handling

- HTTP endpoints :

 - Path : `https://site.com/groups/admins` ? groupId=123
 - Versioning : /v1/ for backward compatibility
 - Routing → clear and versioned (`/groups/v1/admins`
vs action = getAdmins)
 - Use correct HTTP verbs :
 - GET → fetch
 - POST → send data / add
 - DELETE → remove
 - Don't duplicate action in URL (avoid `/getAdmins`)

Common Pitfalls

1) Side Effects

- API should do one thing only
- Avoid hidden operations (e.g. setAdmins also creating grp)
- Break into smaller APIs if multiple actions

2) Atomicity

- Handle dependent operations separately (`creategroup` then `setAdmins`)

Handling Large Responses

- Pagination → client requests in chunks (`offset, limit`)
- Fragmentation → break response into multiple packets (service-to-service)

Data Consistency

- Decide how consistent results need to be:
 - Strong consistency → accurate but slow
 - Eventual consistency → faster; may serve stale data
- Use caching for high-read APIs (comments, profile lookups)
- Graceful degradation: send minimal info (name, id) under heavy load instead of failing.

A good API is

- Intuitive (clear name, right params)
- Precise (only needed data)
- Predictable (no side effects)
- Resilient (handles expected errors + scales with pagination, caching, degradation)

Capacity Estimation

Storage Estimation (YouTube)

Assumptions

- 1B users → 1 in 1000 uploads → 1M uploads/day
- Avg video length ≈ 10 mins, 2 hr → 4 GB → 0.4 GB (low quality)
- Video size: ≈ 3 MB/min → 30 MB/video
 - 400 MB → 2 hr
 - 200 MB → 1 hr
 - 3 MB → 1 min

Raw Storage

$$\text{Total} = 1 \text{M} \times 30 \text{ MB/video} = 30 \text{ TB/day}$$

Redundancy

$$3 \text{ copies} \rightarrow 90 \text{ TB/day}$$

Multi-resolution storage

- Storage in 144p, 240p, 360p, 480p, 720p versions
 - $\frac{1}{16}$
 - $\frac{1}{8}$
 - $\frac{1}{4}$
 - $\frac{1}{2}$
 - $x \text{ MB}$

$$+ \frac{x}{n} \text{ MB}$$

- Storage $\approx 2 \times$ original ($n \text{ MB} + n \text{ MB} = 2n \text{ MB}$)
- Total $\approx 180 \text{ TB} \approx 0.2 \text{ PB/day}$
- Since for 0.4 GB $\rightarrow 0.2 \text{ PB/day}$, 4 GB $\rightarrow 2 \text{ PB/day}$ (high quality)

\rightarrow How high quality \rightarrow low quality (Assumption)
 2 hr \rightarrow 4 GB

optimised code and compression savings = 90%
 2 hr $\rightarrow 4 \text{ GB} / 10 \text{ GB} = 0.4 \text{ GB}$

Metadata Caching Estimate :

- Thumbnails (images, heavier)
- Titles / IDs (lightweight)

Assumption

- Original $\rightarrow 1 \text{ MB}$
- Compressed $\rightarrow 100$ times $\rightarrow 10$ times from both sides
- Display size $\rightarrow 1 \text{ MB} / 100 \approx 10 \text{ KB}$
- Cache only popular videos $\rightarrow 90$ days last uploads
- $10 \text{ KB} \times 90 \text{ days} \times 1 \text{ M uploads/day} \approx 10^9 \text{ KB} \rightarrow 1 \text{ TB RAM}$

Deployment

- 16 GB servers
- 1 TB $\rightarrow 1000 \text{ GB} / 16 \text{ GB} \rightarrow 64$ nodes \rightarrow cache system
- Redundancy (3 copies) + 50% capacity
 $64 \times 3 \times 2 \rightarrow 128 \times 3 \approx 500$ nodes

Video processing Power Estimate :

- \rightarrow convert uploads/day \rightarrow MB/s
- Upload volume $= 10^7 \text{ min/day}$ ($10 \text{ min} \times 1 \text{ M uploads/day}$)
 - Density = 3 MB/min
 - Data / day = $10^7 \times 3$

Video Processing Estimation :

1) Upload volume / day

- $10^7 \text{ min/day} \rightarrow (10 \text{ min} \times 1 \text{ M uploads/day}) \rightarrow 166,667 \text{ hr/day}$

2) Storage

- Assume video size density = 1 GB/hr

$$\text{Raw storage/day} = 166,667 \times 1 \text{ GB} \approx 166 \text{ TB}$$

3) Convert to Throughput

$$\text{Data/day} = 166,667 \times 1024 \approx 2 \times 10^8 \text{ MB/day}$$

$$\text{Seconds/day} = 86,400 \approx 100,000$$

$$\text{Throughput} = 2 \times 10^8 / 10^5 \approx 2000 \text{ MB/s}$$

4) Processing cost / MB

$$\text{Read} = 10 \text{ ms}$$

$$\text{Process} = 20 \text{ ms}$$

$$\text{Write} = 20 \text{ ms}$$

$$\text{Total} = 50 \text{ ms/MB} = 0.05 \text{ s/MB}$$

5) Work / sec

$$2000 \text{ MB/s} \times 0.05 \text{ s/MB} \approx 100 \text{ processes}$$

We need 100 processes in parallel

6) Redundancy + Multiple Resolution (optional)

$$100 \times 3 \text{ (copies)} \times 2 \text{ (resolution)} \approx 600 \text{ processes}$$

Log Structured Merge Tree

Motivation -

- Traditional DBs use B+ Trees

$$\rightarrow \text{Insert} = O(\log n)$$

$$\rightarrow \text{Search} = O(\log n)$$

- Issue : Too many small writes \rightarrow lots of disk I/O + acknowledgments

Goal -

Optimize writes while keeping reads fast enough

Step 1 : Write Optimization

- collect writes in memory → batch them
- Append to a log-like structure (linked list) → O(1) write
- Flush to disk in bulk (reduces I/O + ACK overhead)
- Problem → reads slow → linked list = sequential search → $O(n)$

Step 2 : Sorted String Table (SSTable)

- Before flushing to disk → sort records → chunks
- Persist sorted arrays (SSTables)
- Reads → Binary Search → $O(\log n)$ on each SSTable.

Step 3 : Merge SSTable

- Multiple sorted chunks = multiple binary searches
- Solution → Compaction → merge smaller SSTables into bigger sorted SSTables → similar to merge sort → reduce no. of SSTables to search.

Step 4 : Faster Lookups

- Attach Bloom Filter per SSTable → chunk on step
- Helps decide if a key might exist in an SSTable.
- Avoids unnecessary reads
- Trade-offs → false positive possible → false negatives impossible.

Bloom Filter

1. Bit array of size M (all 0 initially)

2. Insert string → pass through K hash functions → set bits at K positions = 1

3. Query string → check K positions. no false negative possible, false positive possible,

• If any 0 → definitely not present • If all 1 → probably present

Accuracy drops if array too full → reut / rebuild needed

Error probability
 \rightarrow
 $P_{error} = \sum_{i=1}^K P_i$
 \rightarrow
 $P_i = \prod_{j=1}^M (1 - p_j)$
 \rightarrow
 $P_i = e^{-\lambda}$
 \rightarrow
 $P_{error} = e^{-k\lambda}$

Master Slave Architecture

→ Problem

- Typical system : Multiple servers + one database
- Single DB = single point of failure (SPOF)
- If DB crashes → business stops

→ Solution

- Replicate database (copy stored elsewhere, ideally separate hardware).
- Provides backup + fault tolerance.

Types of Replication :

1) Asynchronous

- Master updates , then later copied to slaves
- Lower load on master

Problem → Slave can be out of sync → data loss if master fails

2) Synchronous

- Master update must be copied to slave before commit
- Strong consistency

Problem → Higher latency → extra network round-trip

Master - Slave Roles :

- Master = accepts writes + reads
- Slave = accepts reads only (no writes)
- Writes always go to Master → replicated to slaves

Master - Master Roles (Peer - Peer) :

- Both nodes can accept writes + reads
- Load balancing of writes + fault tolerance
- Split Brain problem
 - If network partition happens , both think they're master
 - Leads to inconsistent data (double detections , etc.)
- Solution

- Add third node (arbiter/consensus node)
- Used for distributed consensus → decides which side is valid
- Ensures only one master state is accepted.

Consensus Protocols

- 2PC (Two Phase Commit) - simple but slow
- 3PC - improvement but still heavy
- MVCC (Multi-Version Concurrency Control) - keeps multiple versions of data, avoids blocking, used in PostgreSQL
- Saga Pattern - long transactions broken into smaller steps with rollback capability.

Why Use Master-Slave?

1. Backup / Fault tolerance : reduces SPOF risk
2. Read scalability : multiple slaves for read-heavy workloads like analytics
3. Eventual consistency acceptable : for non-critical apps
4. Sharding + Master-Slave → scale writes by partitioning data

Location Based Databases

- Used in Google Maps, Swiggy, Zomato, Uber, Dating apps
- Requirements
 1. Distance Measurement → find distance between 2 points
 2. Proximity Search → find all points within X km radius
- Early Approach
 - Assign codes (e.g. ZIP, PIN) to regions
 - Problem → Not accurate → geography like railways, rivers may cause long detours.

1) Better Representation \rightarrow Coordinates :

- Use Latitudes & Longitudes
- Pros
 - \rightarrow Scalable granularity (more decimals \rightarrow more precision)
 - \rightarrow Easy distance computation (Euclidean / Haversine formula)
- Cons
 - \rightarrow Naive proximity search = expensive (check all DB points)

2) Datastructure for Location Search :

→ Quad Trees

- Recursive partitioning of 2D space into quadrants
 - Can adaptively split based on density (e.g. Mumbai \rightarrow deep, Icland \rightarrow shallow)
 - Problem \rightarrow range queries in 2D are expensive
- Solution : Reduce 2D to 1D
- Range queries are efficient in 1D (segment trees, interval trees).
 - Convert 2D spatial data into 1D line while preserving proximity

3) Fractal curves for 2D \rightarrow 1D Mapping :

- Z-curve: simple, but poor locality preservation
- Hilber curve:
 - Space filling curve mapping $2D \rightarrow 1D$
 - Preserves locality well (nearby points in 2D \rightarrow close on 1D line)
 - Recursive \rightarrow can zoom infinitely for higher precision
- Used for efficient range + proximity range queries.

1) Proximity Queries with Hilbert Curve

- Map 2D points → positions on 1D Hilbert curve
- To find nearby point:
 - Take a \pm threshold around a position
 - Eg. pos(29), threshold $\pm 6 \rightarrow$ range [23, 34] → nearby region
- Efficient for approximate proximity search
- Edge cases:
 - Sometimes far points appear close (false positives)
 - Sometimes close points appear far (false negatives)
 - Deeper recursion reduces errors

Consistency in Distributed Systems

- Consistency = multiple copies of data should match each other
- If update in one place, all replicas must eventually or immediately reflect the same value

1) Single Server

- Easy consistency → only one copy of data
- Problems:
 - Single point of failure (if server dies → data lost)
 - Vertical scaling limit (can't scale beyond supercomputer)
 - High latency (users far away get slow responses)

2) Multiple servers without Sharing

- Servers host only local data → reduced load per server
- Problems:
 - No global data sharing
 - Latency if cross region requests
 - Still single point of failure per region.

3) Replication (Multiple Copies of Data)

- Store same data in multiple regions →
 - No single point of failure
 - Reduced latency (read / write locally)
- Problem / Challenge → keeping replicas consistent
 - Update Propagation : updating one copy means syncing all others
 - Message Delivery failures : updating or acknowledgement may get lost
- Two general problems
 - can't guarantee both sides know for sure the update succeeded
 - Infinite ack-of-ack loop problem.

Solutions & Protocols

→ Leader - follower Model

- One server = leader (handles all writes)
- Other servers = followers (replicate data from leader)
- Simplifies consistency management
- Problem :
 - If leader fails, need re-election
 - Reads may see stale data until replication finishes

→ Consistency vs Availability Tradeoff

- If you wait for replication → strong consistency, low availability
- If you respond immediately → high availability, weaker consistency
- This is CAP Theorem in action

→ Two Phase Commit (2PC)

- Used for distributed transactions across multiple systems
- Steps :-

1. Prepare Phase → leader asks followers to "prepare" → log updates but don't commit yet
2. Commit Phase → if all acks, leader sends "commit"

Else, rollback.

- Ensures atomic commit
- Problem → slow + blocking (if a follower hangs, system waits)

Types of Consistency Models

- Strong Consistency → Every read sees the latest write
- Eventual Consistency → Updates propagate asynchronously, replicas eventually become consistent
- Causal Consistency → Related operations preserve order, but not all operations are globally ordered.

Consistency Tradeoffs → ^(high)Latency and Availability (low)

ACID (Properties of transaction in database)

A - Atomicity → All or nothing

C - Consistency → Preserving database invariants

I - Isolation → concurrent transactions are isolated from each other

D - Durability → data is persisted after transaction is committed even if a system failure

CAP (Theorem for distributed system)

C - Consistency → every read gets the latest write (all nodes show same data)

A - Availability → every request receives a response (no downtime)

P - Partition tolerance → system continues to operate even if network partitions occur.

• In distributed system, you can only guarantee 2 out of 3 →

Asynchronous API ?

Async is core to scalable systems. Perform tasks without waiting for others to finish.

Frontend need for

- Handling user interactions smoothly
- Running background tasks (API calls, animations, data fetches) without freezing UI

Backend need for

- concurrent requests hitting server
- Parallel execution (querying DB + calling on API + writing logs simultaneously)

why Async important ?

- Higher efficiency & lower cost
- Better customer experience → no waiting
- Real world impact
 - Amazon → +100ms latency = +1% revenue
 - Google → +500ms delay = -20% user

Types of Async

1. Concurrency

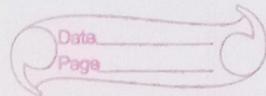
- Single worker switching between tasks
- One task at a time, but rapid switching

2. Parallelism

- Multiple workers doing tasks simultaneously
- Faster, but higher cost of idle resources possible

3. Mix of Both

- Most real systems combine concurrency + parallelism
- Maximum throughput + low latency.



Drawbacks

- Too much concurrency → context switching → thrashing (wasted CPU cycles)
- Too much parallelism → idle / wasted resources → high cost
- Complexity → harder to write / debug code → race conditions
→ requires language features to manage.

OAuth (Open Authorization) :

OAuth ≠ Authentication (who you are)

OAuth = Authorization (what you can access)

But in practice, used for authentication too

Problem with Passwords

- Too many websites → too many passwords
- Pain users (memory / managers)
- Pain for startups (drop off during signup)

OAuth flow :

1. User signs up → instead of creating password, site say "Login via Google / Facebook / etc".
2. Redirect to provider (Google)
3. Provider checks if user exists → asks user for permission to share limited data (e.g. name, email)
4. Provider sends back verified user details + access token
5. Website uses token for session
6. User can revoke access anytime (token invalidated from Google dashboard).

War Story (Startup Example)

- Before OAuth : ~50 signups / day
- After OAuth : 2X signups immediately
- Biggest impact : mobile users (hate typing passwords)

Advantages :

1. Outsourced security → no need to handle passwords
2. Standardized protocol → simpler implementation
3. Easy onboarding → click & sign in
4. Mobile friendly → frictionless signup
5. Email verification handled by provider
6. can integrate multiple provider (google, github, facebook, etc)

Drawback :

1. Dependency on provider → if Google / Meta is down, your signups fail
2. Less flexibility → limited to what provider allows
3. Risk of data sharing → users may see competitor ads
4. Account deletion issue → if provider account deleted, user may lose access
5. Vendor lock in → provider can ban you.

Security Concepts

- Authentication = verifying identity
- Authorization = granting permission to access resources
- Other mechanisms
 - Token based authentication
 - Single sign On
 - Access Control Lists & Rule Engines
 - DDoS protection, secure CDN / video access, DRM for content.

Virtualization & Containers

Problem :

- Capacity planning → gives compute, memory, storage needed → buy hardware
- Issues
 - If too small → scale problem
 - If too big → high upfront cost & wasted resources
 - Scaling horizontally → buying more computers
- Resources
 - Multiple apps/users share same machine → resource contention
 - OS manages CPU, Memory, IO, Disk allocation
 - Still → one program can hog resources → affect others.

Virtual Machines (VMs)

- VM = OS running on top of OS (via hypervisor)
- Provides strong isolation → each VM gets dedicated resources
- Benefits
 1. Cloud computing model - rent VMs instead of buying servers
 2. Platform independence - e.g. run Linux on Windows
 3. Dynamic provisioning - start/stop VMs easily
- Drawbacks
 - slow boot times (like starting full computer)
 - Heavy resource usage

Containers

- Lightweight virtualization → isolates apps + dependencies, not full OS
- Faster boot / teardown vs VMs (no full OS)
- Mount / unmount process → like attaching file system
- Docker → made containerization simple & developer friendly
 - Developers specify : resources + OS
 - Docker handles hardware / VM abstraction

- Advantages

- Fast provisioning (boot time)
- Easy portability (code + environment)
- Efficient resource usage compared to VMs

- Drawbacks

- Slightly slower than native execution
- Possible firewall / networking issues
- Overhead of container management not worth it for simple apps.

Service Discovery & Health Checks

- Server side focus = reliability & availability, not raw efficiency
- User trust and user services more if they are always available → leads to revenue.

Health Checks

- Health Service monitors all servers (s_1, s_2, s_3, \dots)
- Regularly sends heartbeat messages ("Are you alive")
- If server doesn't respond twice → marked dead / critical
- Dead server → remove from load balancer + restart on monitor another machine.

Problem → sometimes OS responds "Yes" but application is dead

Solution → Two way Heartbeat:

1. Health service → Server ping

2. Server → proactively sends "I'm alive"

Helps avoid zombies (machines running but unusable)

Service Discovery

• New service runs on multiple servers

• Must register with Load Balancer

- IP addresses + ports of active instances

- LB maintains a snapshot of all services
- Other services don't need to track IPs → always query LB
- Snapshots are versioned and persisted (DB)
- LB can cache service info for efficiency.

Health checks + Service discovery Integration

- When snapshot changes (new/dead instance)
 - Health Service verifies new servers are alive
 - Keeps connections to service ports.
- Ensures dynamic system awareness (what's alive, what's dead, where requests should go).

Rate Limiting ☺

Problem : Thundering Herd

→ Sudden spike in requests = crushes servers

→ Example : 4 servers, one crashes → load redistributes → others overload & crash (cascading failure)

→ Cascading Failure : chain reaction of server crashes

Solutions

1) Rate Limiting (Queues + Capacity)

- Each server has queue + max capacity (e.g. 300 req/sec)
- Extra requests → dropped (fail response)
- Client retries later → avoids overload
- Return temporary vs permanent errors
 - Temporary : retry later (server busy, PB slow)
 - Permanent : bad request (don't retry)

2) Scaling for Events

- Pre scaling → add servers before known spikes
- Auto scaling → cloud auto adds/removes servers
- Viral traffic → fallbacks to rate limiting

3) Job Scheduling

- Avoid running huge jobs at once (e.g. 1M new yr mails)
- Break into batches (1K/min) → manageable load
- Batch processing improves reliability

4) Popular Posts Handling

- If a celebrity posts → millions of notifications
- Use batch processing + jitter (staggered sending)
- Metadata (likes / views / comments) → serve approximate stats instead of exact → saves DB load.

5) Good Practices

- Caching → store frequent responses, reduce DB queries
- Gradual Deployment → release updates to small % servers first
- Coupling (with caution)
 - Cache non critical data (profile pic) locally to reduce external calls.
 - Avoid caching sensitive data (auth, token, password)

Anomaly Detection in Server Systems

- Monitor system health
- Detect anomalies (unusual patterns)
- Alert engineers early.

Design Principles

- Prefer false positives > false negatives
 - cheaper for engineer to check false alarm than to miss a real issue
- Uniform & standardized metric collection from all services
- Use sidecar (Service Mesh)
 - Handles routing, messaging, logging
 - Collects metrics → sends to anomaly detection engine

Detecting Anomalies

1. Threshold based

- Define upper/lower limits (static or from historical data)
- Simple but error prone (ignores long term trends)

2. Dynamic Filters

- Low pass & High pass filters → follow graph trend
- Sharp deviations = anomaly
- Limitation: ignores seasonal/temporal effects (e.g. Christmas spikes)

3. Seasonality Awareness

- Compare with historical seasonal data (e.g. last yr's December)
- Prevents false alarms during predictable peaks.

4. Isolation trees (ML)

- Based on decision trees
- Partition data → anomalies = points isolated with fewest cuts
- Works well for time series anomaly detection.